



Evaluating Shared Memory Heterogeneous Systems Using Traverse-compute Workloads

Yanwen Xu

University of California, Santa Cruz

Ang Li

Princeton University

Tyler Sorensen

University of California, Santa Cruz



Highlights

Many applications in edge computing can benefit from utilizing tree data structures to accelerate their workloads

Showed how open-source hardware can be leveraged to accelerate a specific class of tree algorithms, which we call *traverse-compute*

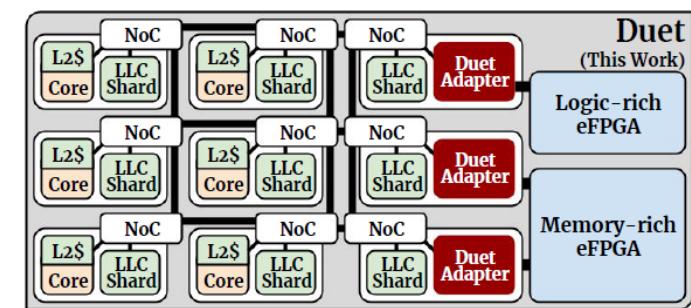
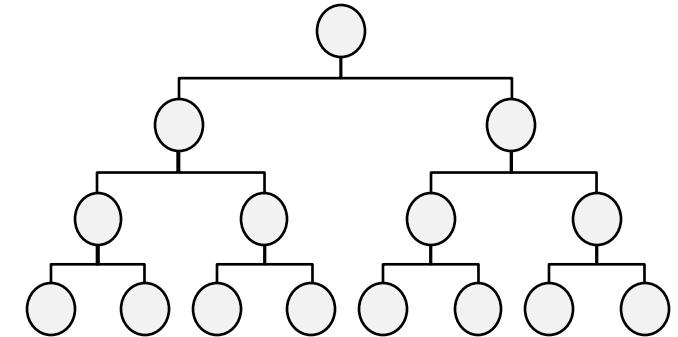
Evaluated open-source heterogeneous architecture called **Duet**, using a recently published open-source framework and benchmark suite **Redwood** and **Grove**

- w/ 9 pragmatic traverse-compute applications

Achieved

- **13.53x** highest speedup
- **6.43x** geomean speedup

Insight: Traverse-compute workload has natural heterogeneous decompositions on modern shared memory system-on-chips



Motivation: Accelerating Computations at Edge

Edge computing are getting popular ...

But they has **constraints**

- *e.g., energy or latency requirement*

Application of edge computing

- *Surveillance cameras*
- *Autonomous vehicles*
- *Mobile gaming*



Motivation: Accelerating Computations at Edge

Edge computing are getting popular ...

But they has **constraints**

- *e.g., energy or latency requirement*

Application of edge computing

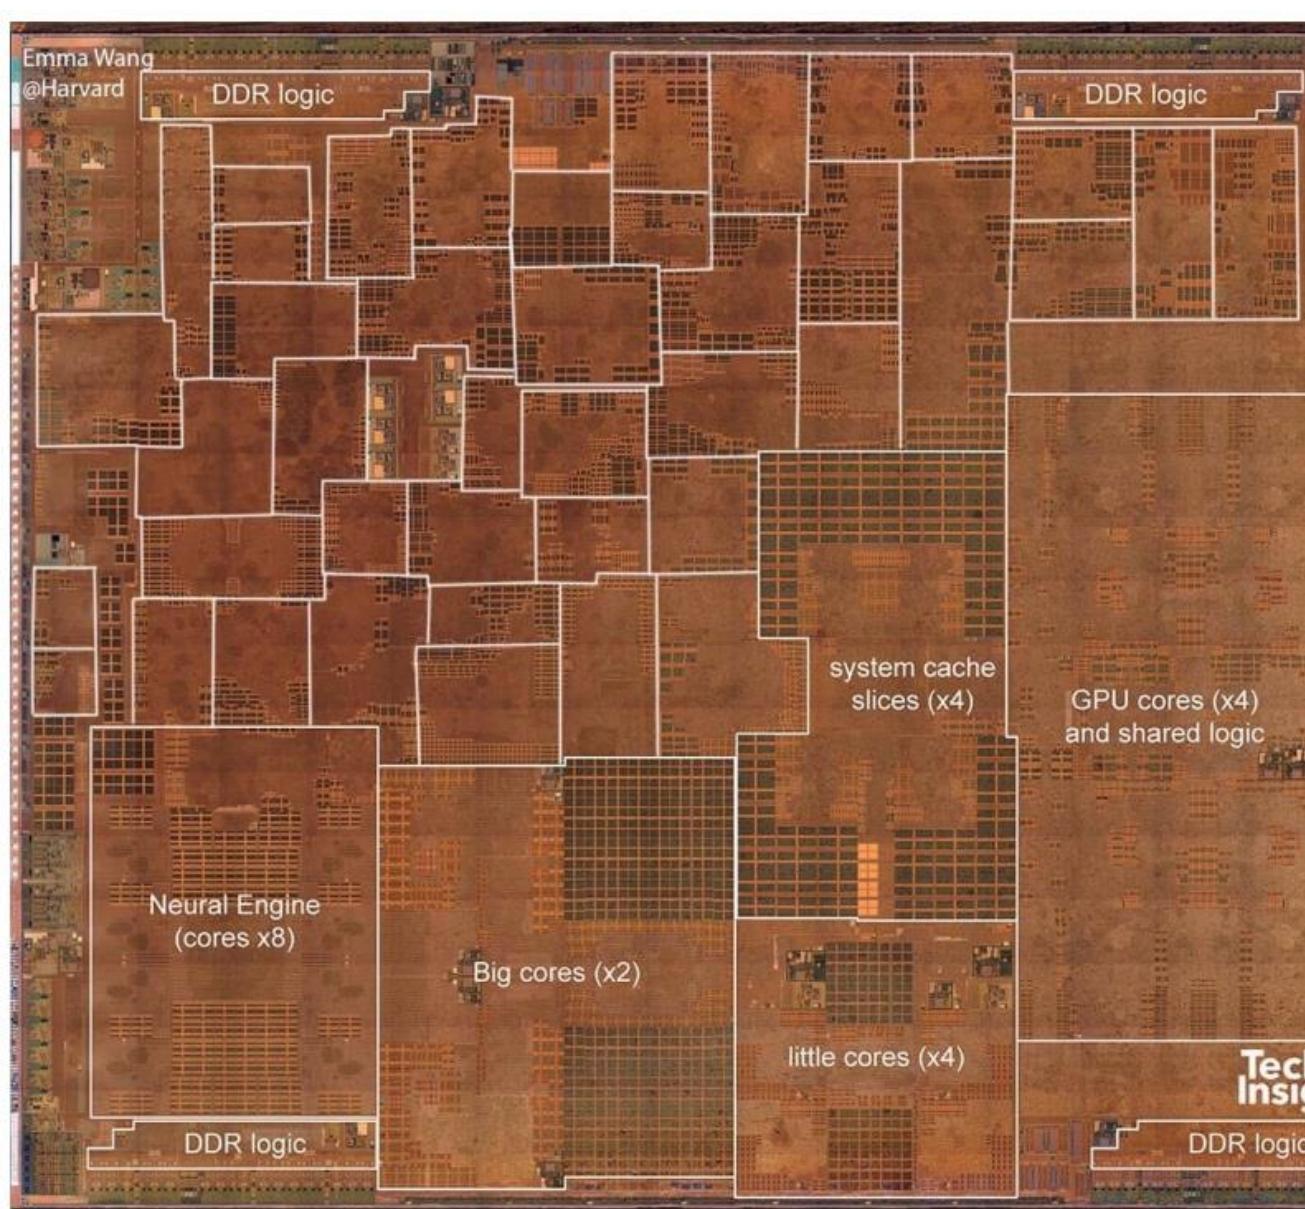
- *Surveillance cameras*
- *Autonomous vehicles*
- *Mobile gaming*

Modern edge devices are becoming increasingly heterogeneous

- w/ specialized *Processing Units* (PUs)



We need to efficiently utilize these available system resources

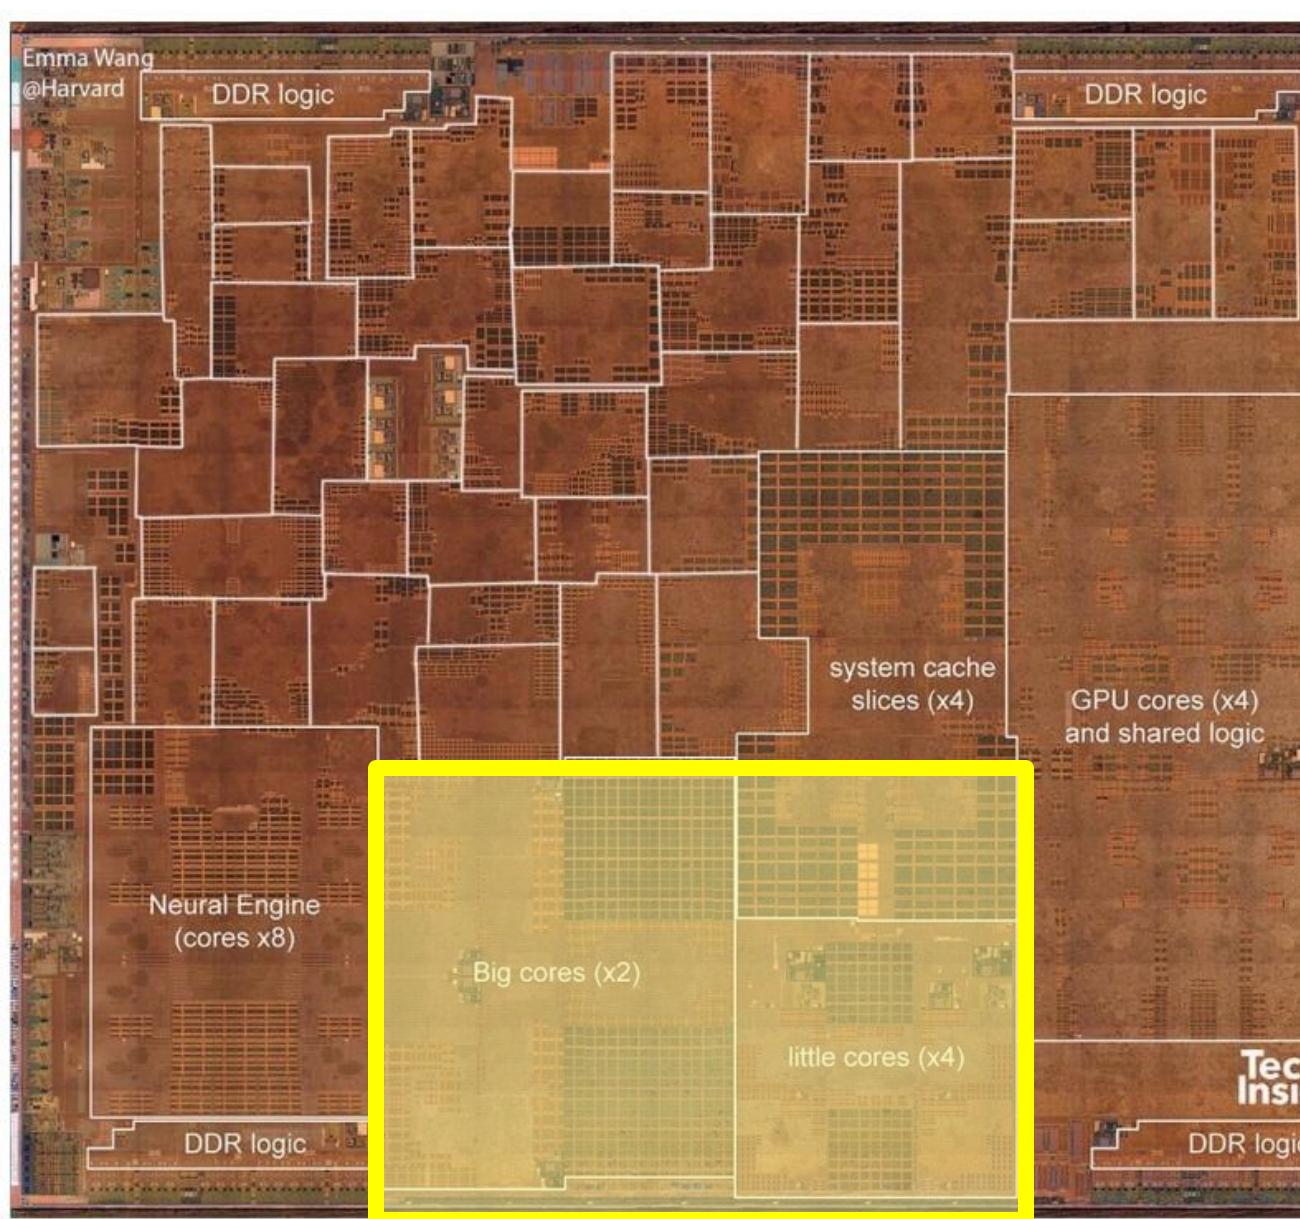


From David Brooks lab at Harvard:

<https://vlsiarch.eecs.harvard.edu/research/accelerators/die-photo-analysis>

What do we mean by resources?



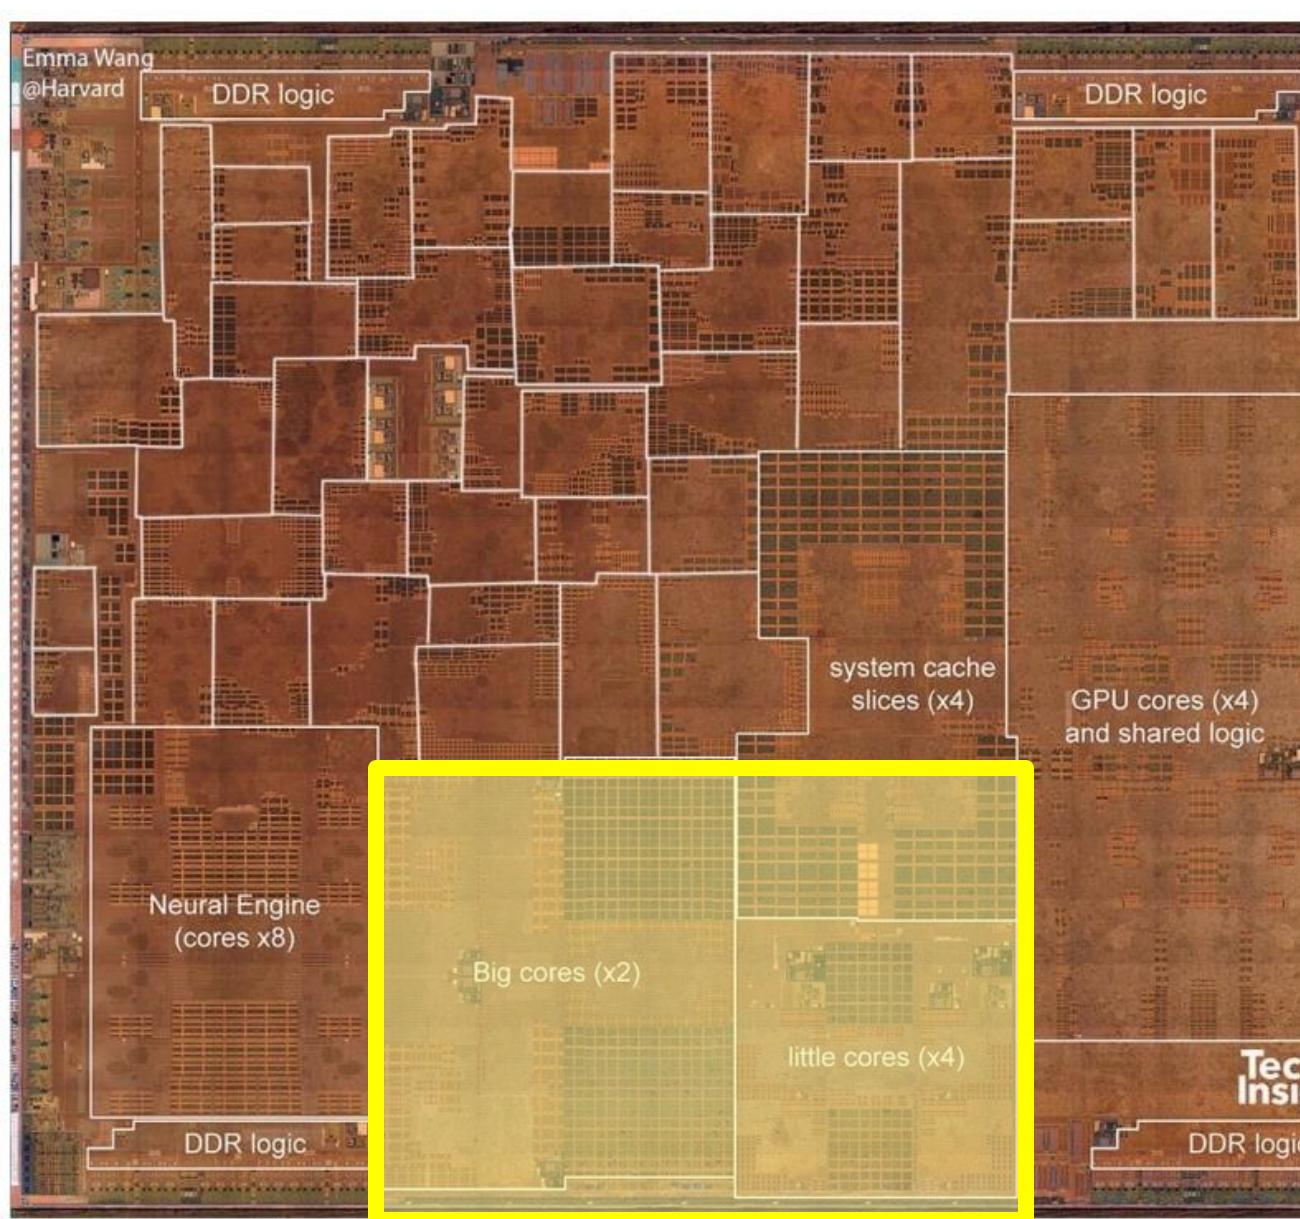


From David Brooks lab at Harvard:
<https://vlsiarch.eecs.harvard.edu/research/accelerators/die-photo-analysis>

What do we mean by resources?

- E.g., less than **20%** of the die area of an iPhone contains the CPU
- The rest contains specialized *Programmable Accelerating PUs (PAPU)*
 - e.g., integrated GPUs, FPGAs
 - Interconnected to a shared memory hierarchy
 - *Shared Memory Heterogeneous System (SMHS)* enables efficient communication between PUs





From David Brooks lab at Harvard:
<https://vlsiarch.eecs.harvard.edu/research>

What do we mean by resources?

- E.g., less than **20%** of the die area of an iPhone contains the CPU
- The rest contains specialized *Programmable Accelerating PUs (PAPU)*
 - e.g., integrated GPUs, FPGAs
 - Interconnected to a shared memory hierarchy
 - *Shared Memory Heterogeneous System (SMHS)* enables efficient communication between

How can workloads efficiently utilize each PU?



Processing Units (PU) Characteristics

CPU

Features: High-performance cores, reorder buffer, load store queue, ...

- + Latency optimized
- Limited throughput

Good for **irregular** programs

Programmable Accelerating PUs (PAPU)

GPU

Features: SIMT (Single Instruction, Multiple Threads) execution, coalesced memory access

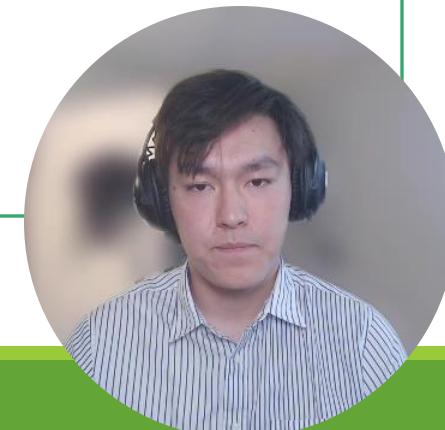
- + Throughput optimized
- Warp Divergence

Good for accelerating
compute-intense programs

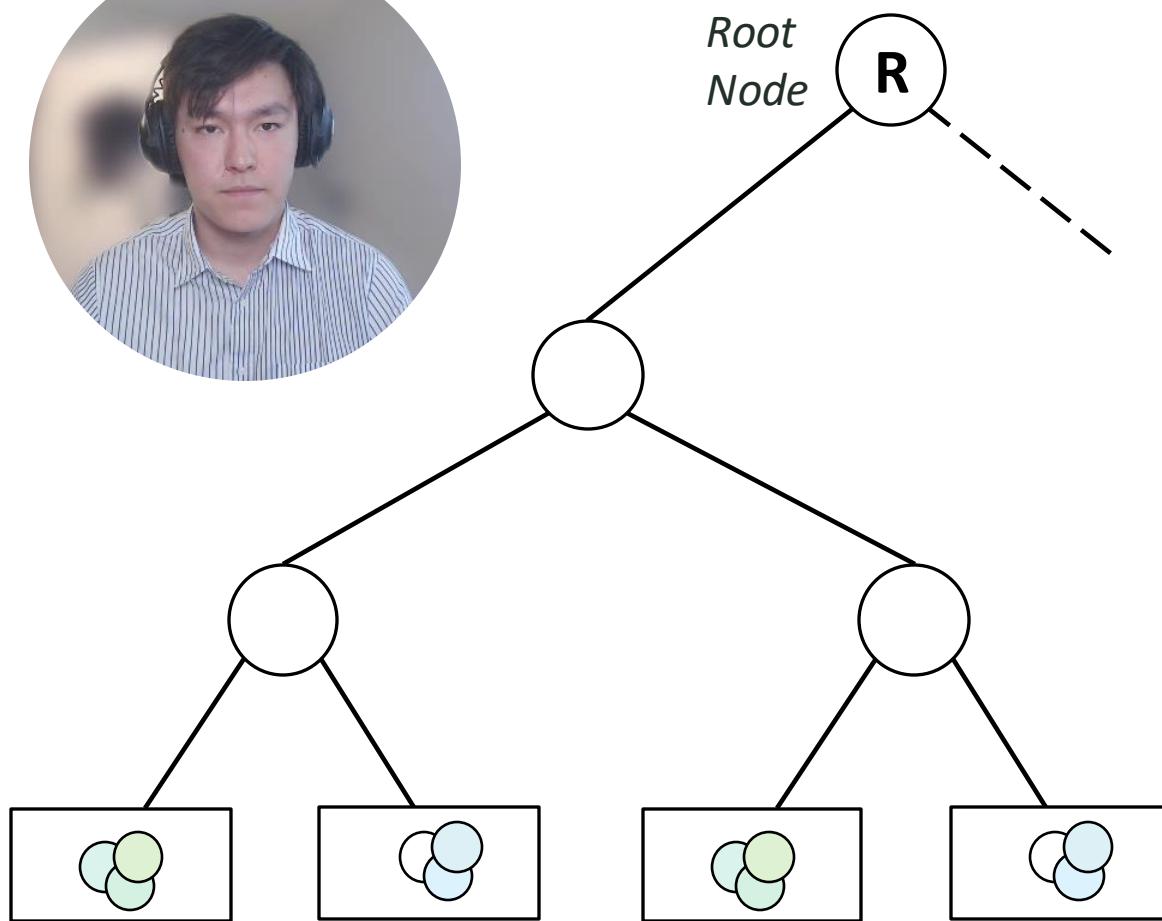
FPGA

Features: Specialized tasks, Pipeline parallelism

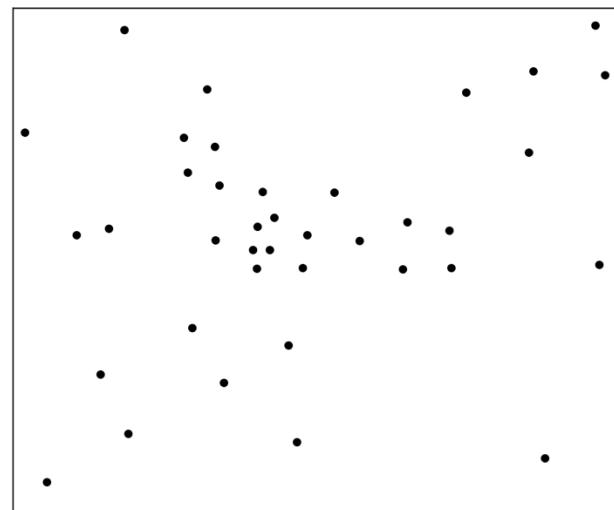
- + Close to ASIC performance
- Orders-of-magnitude harder to program



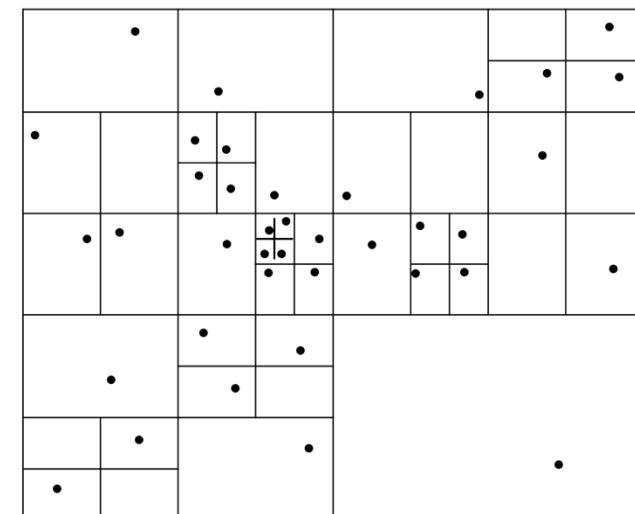
Trees on the edge



- Edge applications need to process a large amount of data
- They can utilize **tree structures** and traversals to perform edge tasks
 - E.g., *octree, k-dimensional tree*

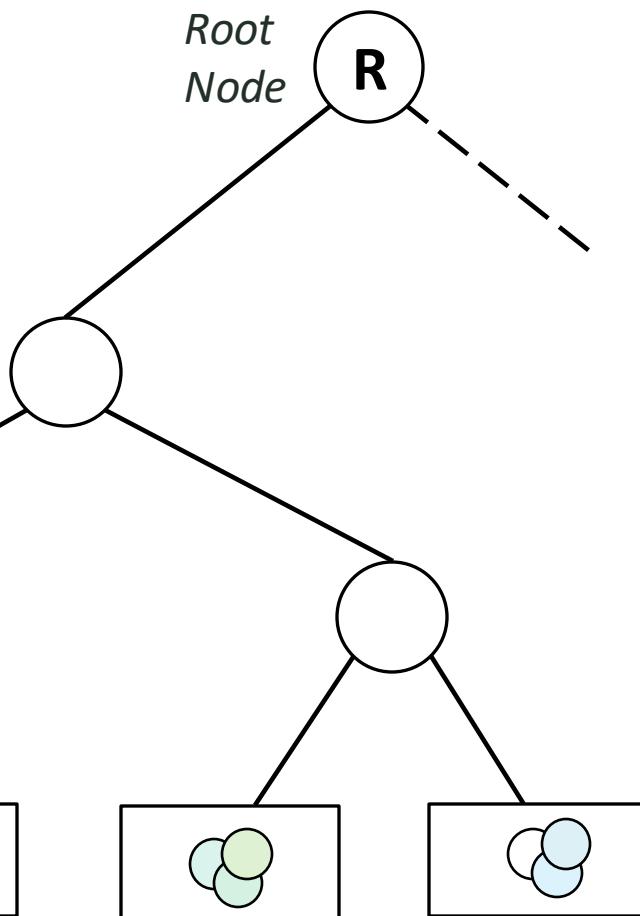


Input data

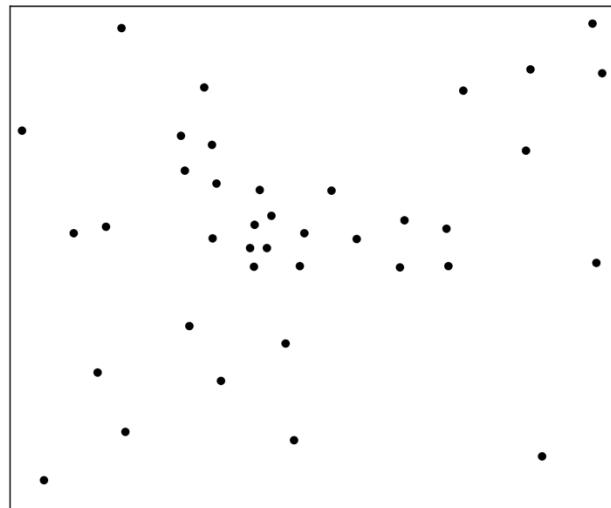


Spatial Partition

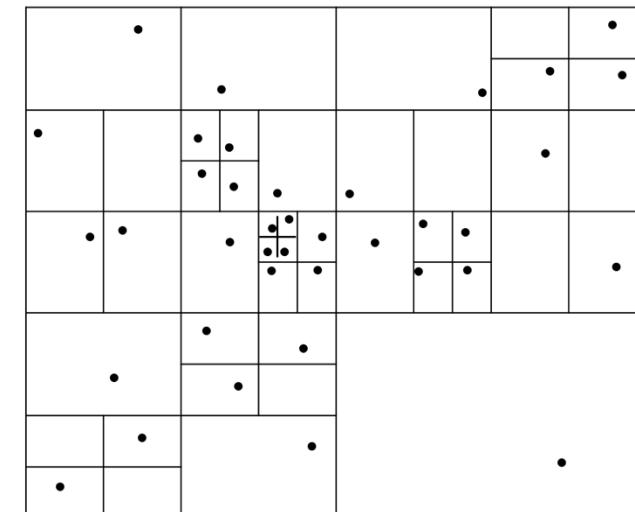
Trees on the edge



- Edge applications need to process a large amount of data
- They can utilize **tree structures** and traversals to perform edge tasks
 - E.g., *octree, k-dimensional tree*
- The dataset are organized into a hierarchical tree structure, allowing data to be **efficiently** searched from $O(n)$ to $O(\log n)$

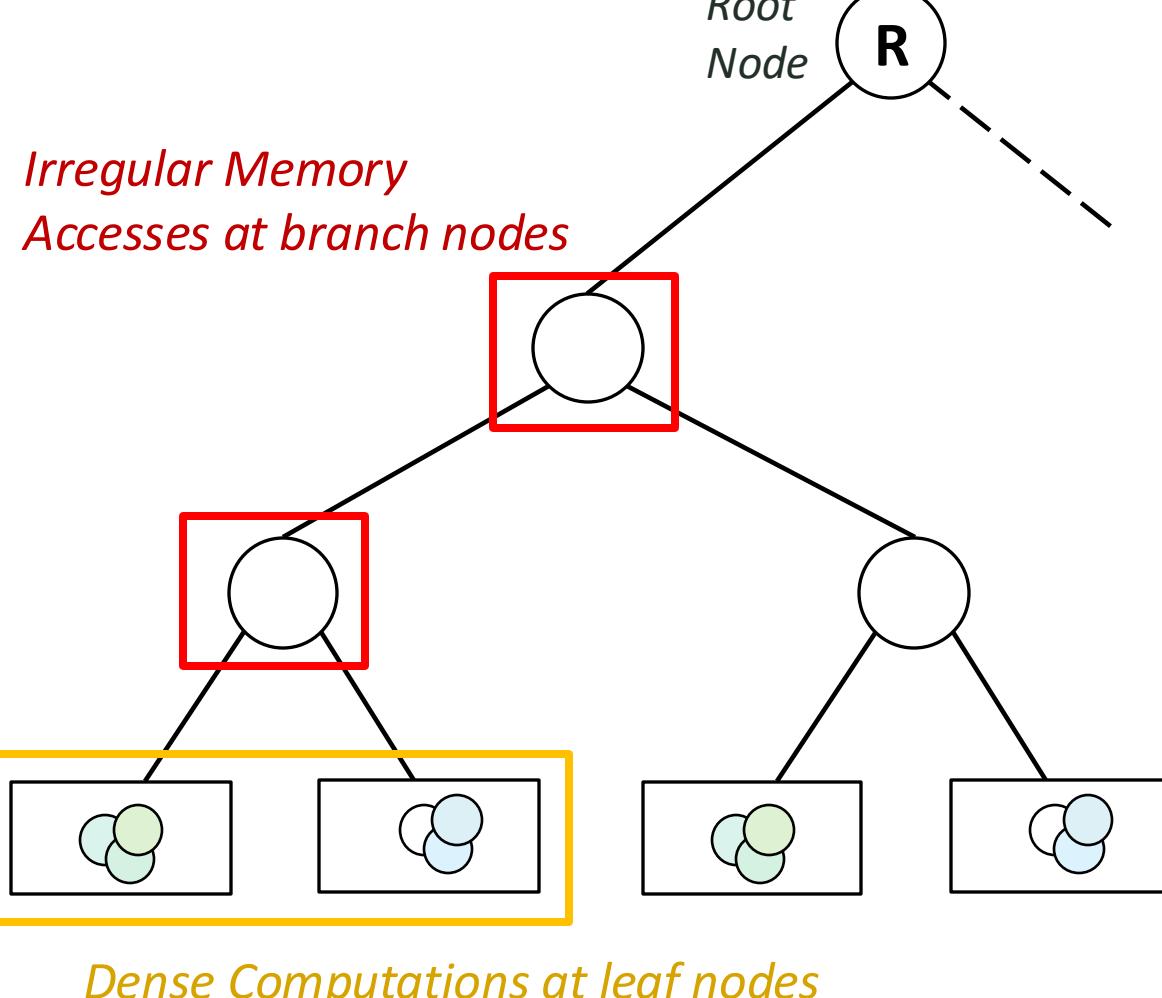


Input data



Spatial Partition

Traverse-Compute Workloads



- Repeatedly traversing a sparse tree structure
- Each traversal consists of
 - Indirect memory loads at branch nodes (**Red box**)
 - Dense data to be processed at leaf nodes visited (**Orange box**)
 - Computing pairwise interactions (e.g., Euclidean distance)
 - Reductions (e.g., sum, min)
- Example workloads:
 - Barnes-hut Algorithm (octree)
 - Nearest Neighbor (kd tree)
 - Ray Tracing (BVH)



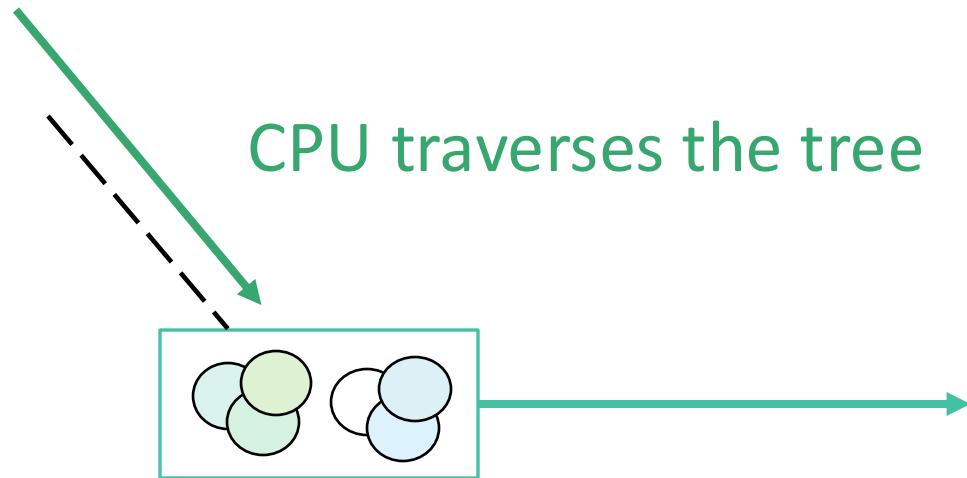
Decomposing Traverse Compute Workloads

Tree applications can benefit from fine-grained acceleration

CPUs are good at handling dynamic control flows and tolerating indirect memory loads

PAPUs are good at accelerating regular, compute-intense operations

A natural heterogeneous approach is to



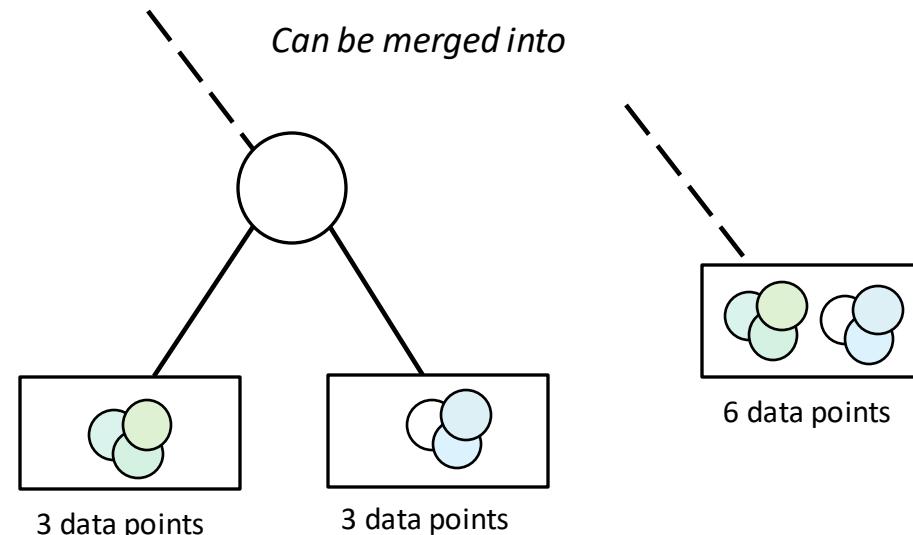
CPU traverses the tree

Computations at leaf nodes are offloaded to the accelerator



Accelerating Traverse-compute workloads on SMHSs

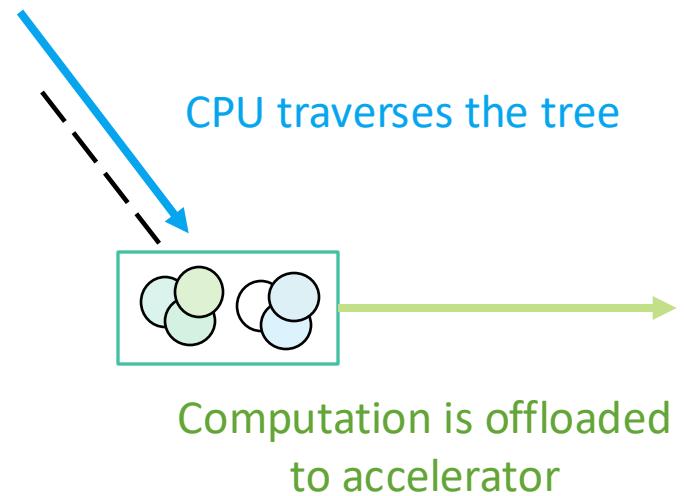
- The tree can be **parameterized** by how many data points exist on the leaf nodes.



Larger node sizes tradeoffs

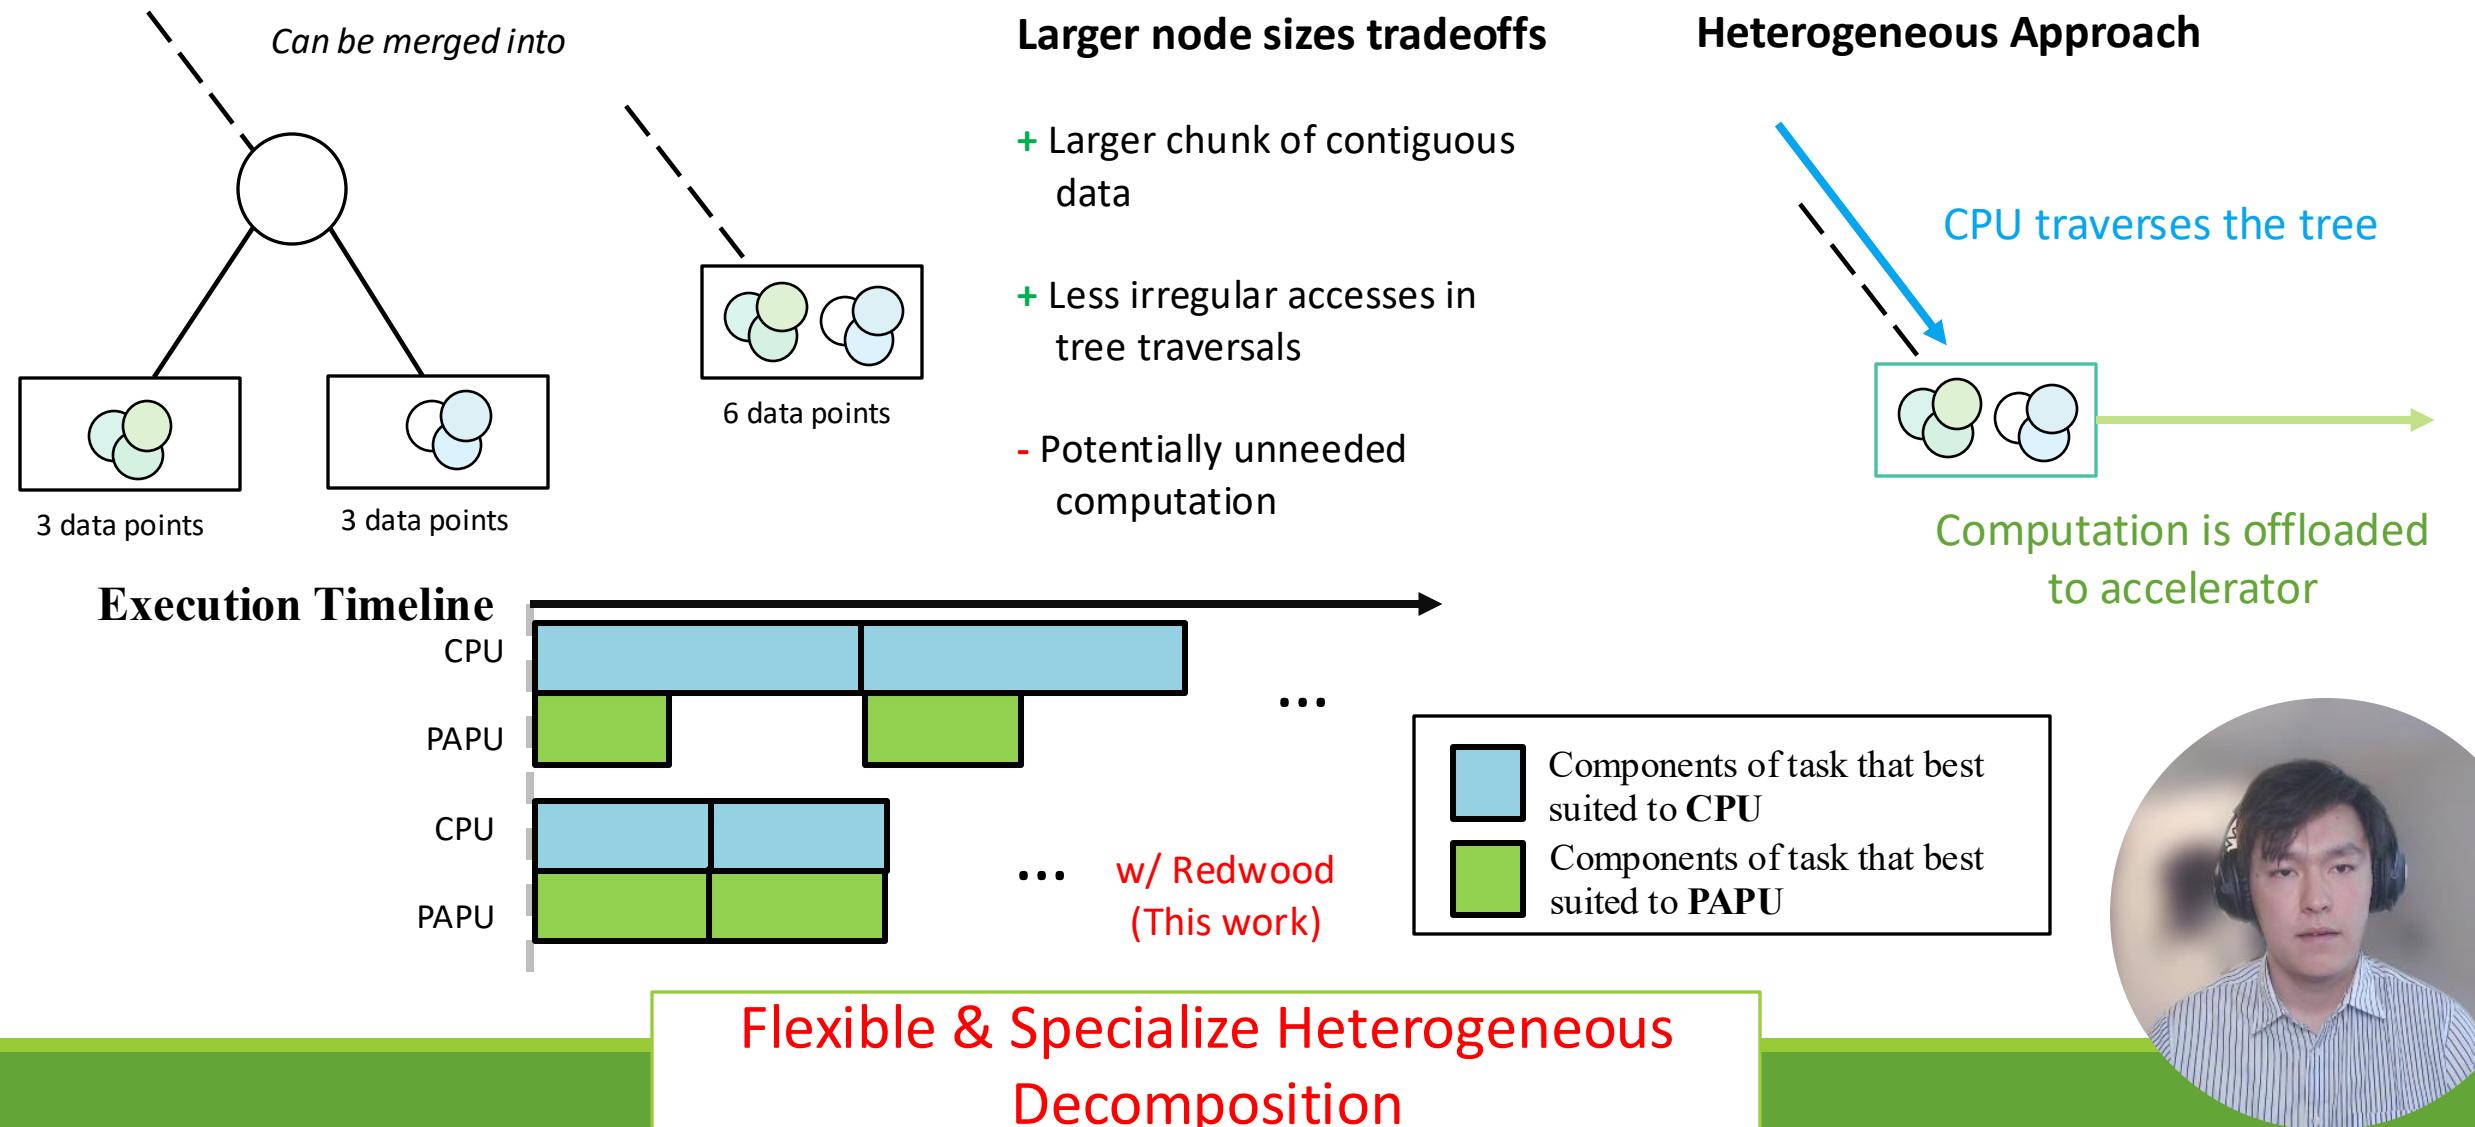
- + Larger chunk of contiguous data
- + Less irregular accesses in tree traversals
- Potentially unneeded computation

Heterogeneous Approach

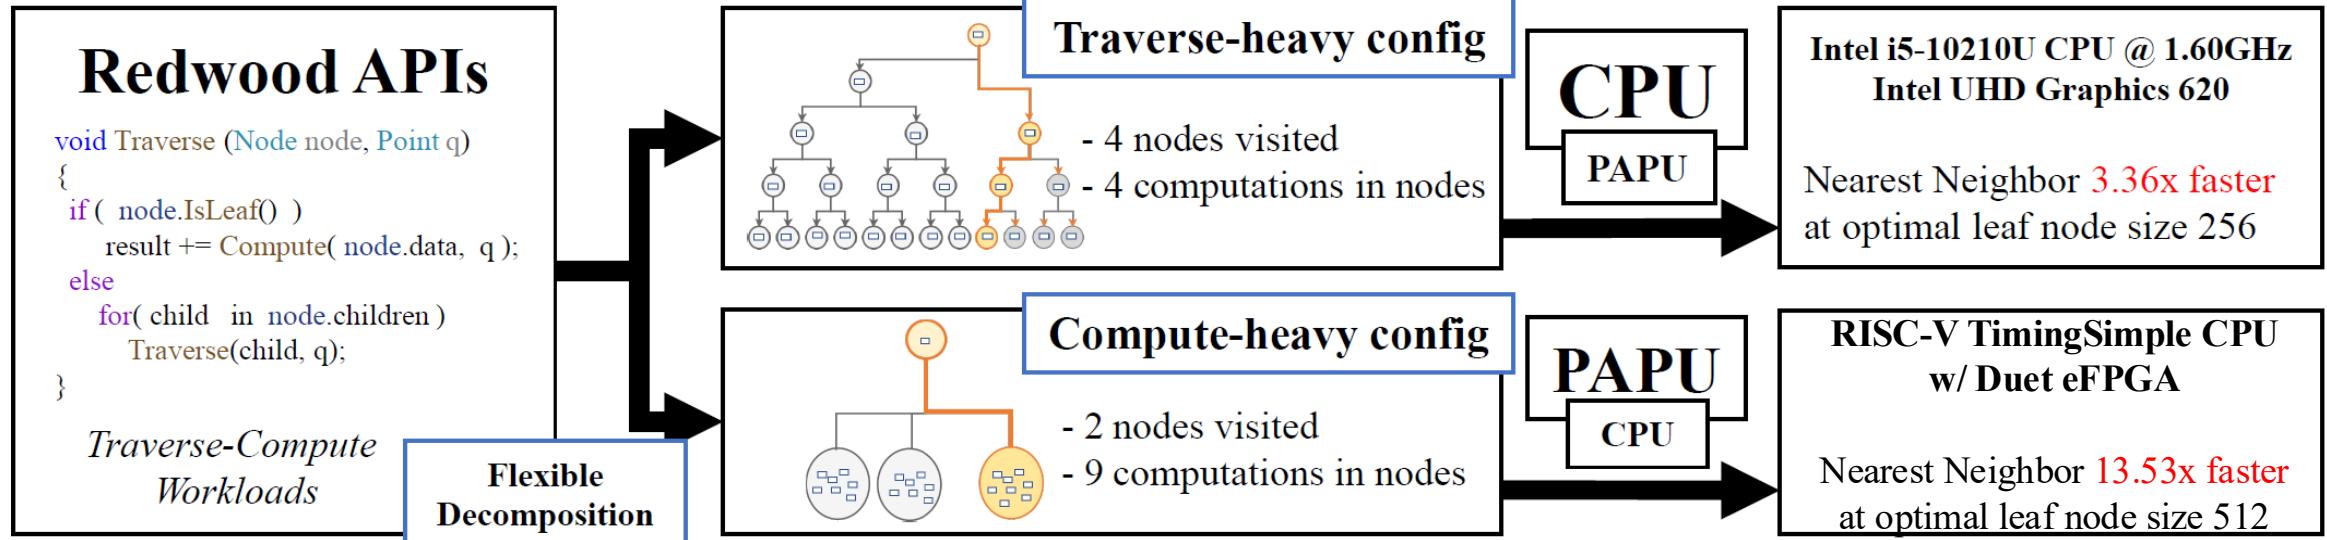


Accelerating Traverse-compute workloads on SMHSs

- The tree can be **parameterized** by how many data points exist on the leaf nodes.



This work: Redwood Overview



Users implement tree applications using our APIs



KNN based Facial recognition

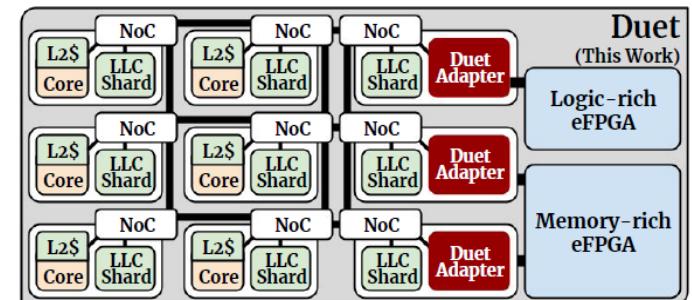
Target systems w/ different CPU/PAPU throughputs



Intel SoCs



Nvidia SoCs



Duet



Redwood: APIs and Data Structures

CPU Sequential Code (NN)

```
tree = KDTree()  
min_dist = 99999.9999  
def traverse(node, q):  
    if is_leaf(node):  
  
        # Reduce Leaf Node  
        for i in range(node.leaf_size):  
            kernel_func(q, node.data[i])  
  
    else:  
        dist = compute_dist(q, node.data[0])  
        min_dist = min(min_dist, dist)  
        traverse(node.leaf_child)  
        if check_other_side(dist):  
            traverse(node.right_child)
```

*Implemented
using:*

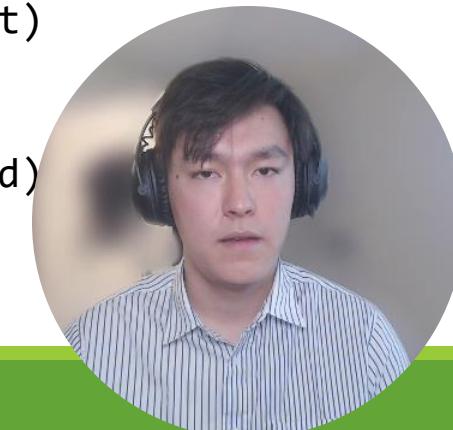
w/ Redwood API

```
tree = KDTree(leaf_size=32)  
redwood_set_query(q)
```

```
def traverse(node, q):  
  
    if is_leaf(node):  
        redwood_compute_leaf(node.data())
```



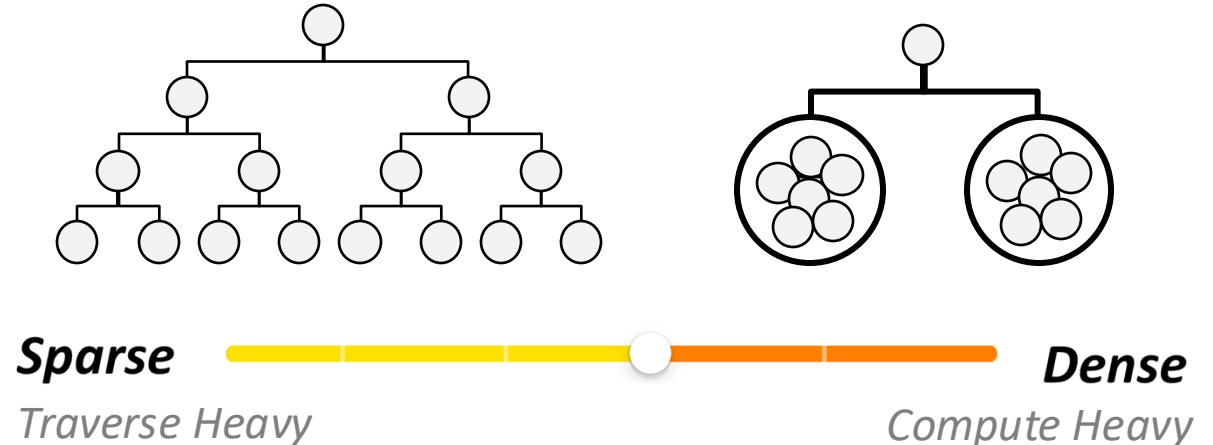
```
else:  
    dist = compute_dist(q, node.data[0])  
    min_dist = min(min_dist, dist)  
    traverse(node.leaf_child)  
    if check_other_side(dist):  
        traverse(node.right_child)
```



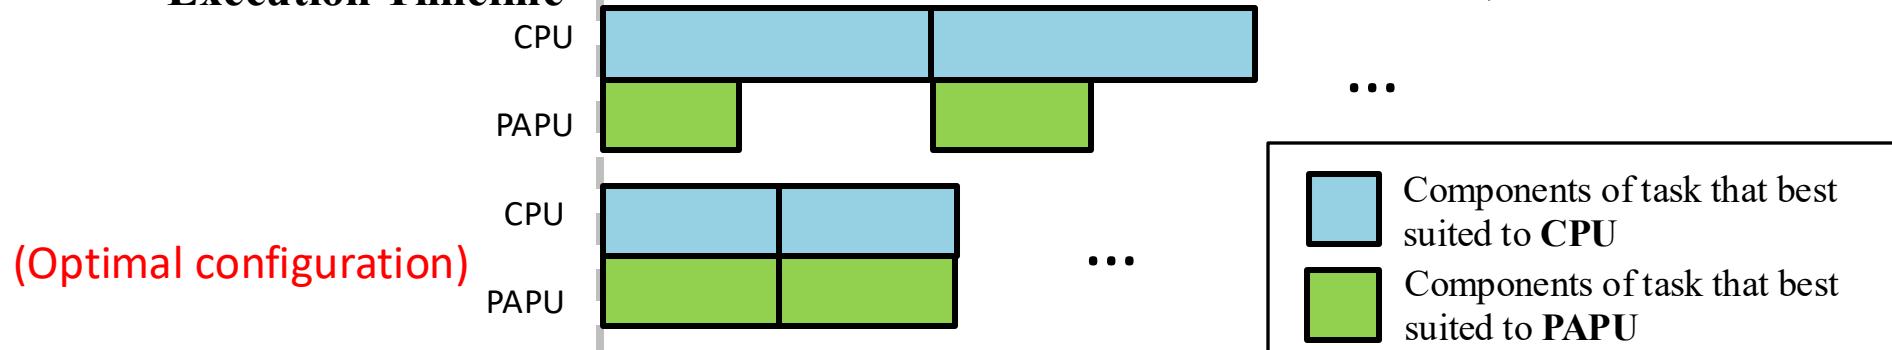
Redwood Heterogenous Optimizations

Flexible Leaf Size Configuration

- Adapt to various heterogeneous systems with different relative throughput between the CPU and the PAPU



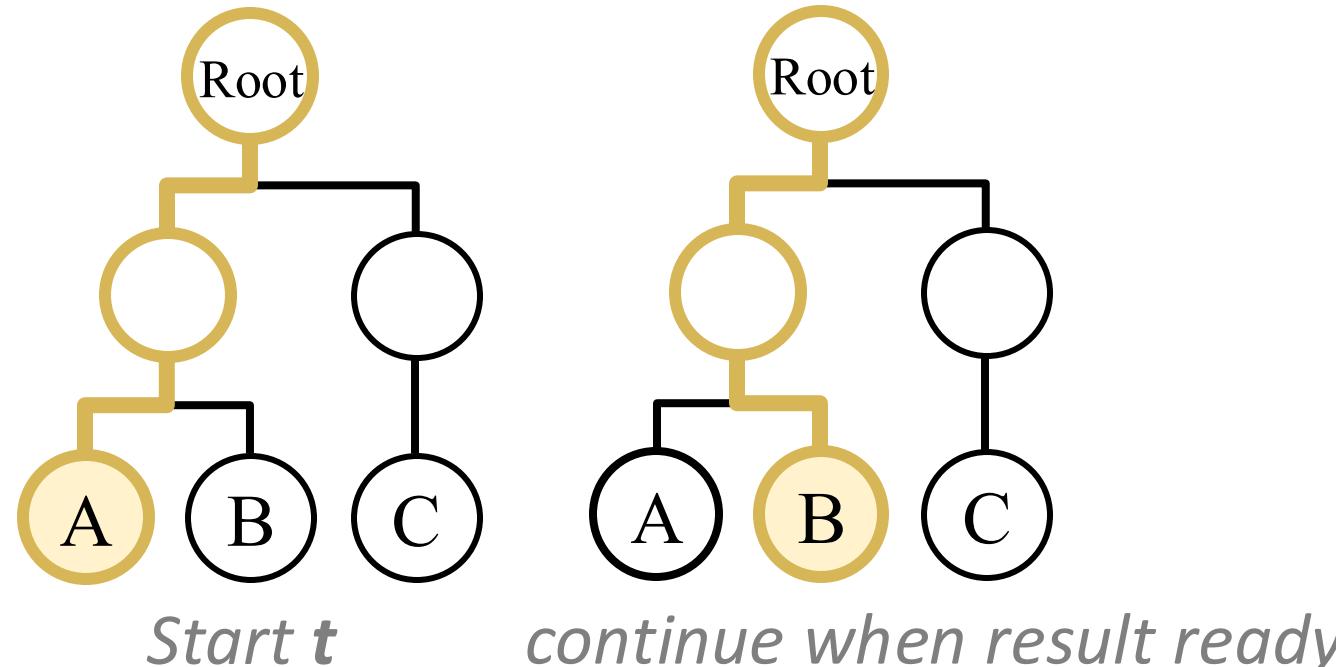
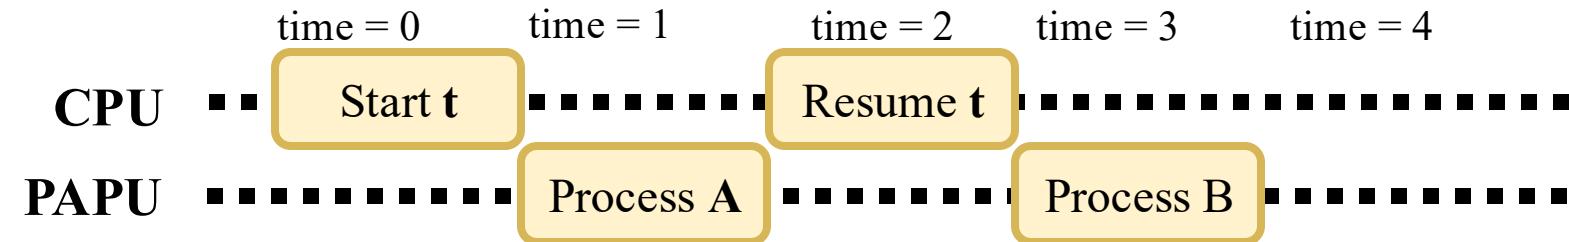
Execution Timeline



Redwood Heterogenous Optimizations

Traverser Runtime

- Allow a single CPU thread to execute many traversals concurrently to **avoid stalling** when a traversal depends on a PAPU accelerated value

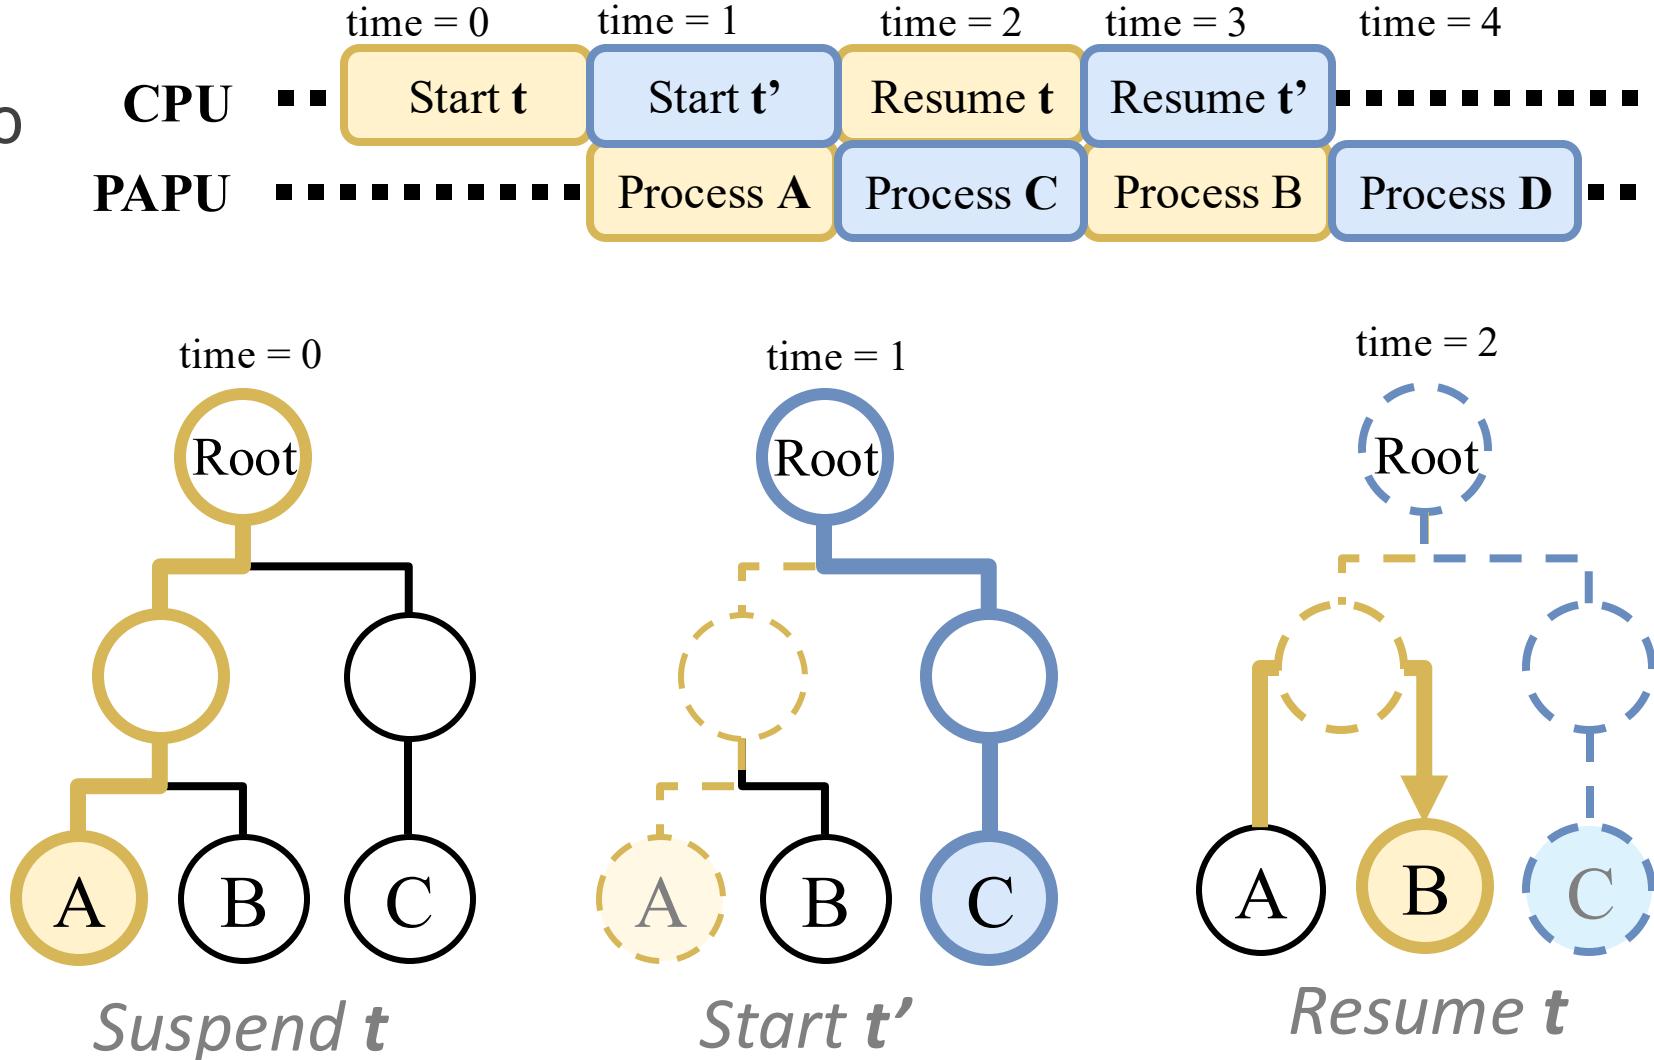


Redwood Heterogenous Optimizations

Traverser Runtime

- Allow a single CPU thread to execute many traversals concurrently to **avoid stalling** when a traversal depends on a PAPU accelerated value

- Lightweight Coroutine*
 - Suspend*
 - Resume*



Grove: Benchmark Suite for SMHS

Grove contains **9** traverse-compute workloads

Can be found in many applications

- *Astrophysics*
- *Facial recognition*
- *Anomaly detection*
- *Outlier detection*
- *Particle simulation*

Tree Structures

- Octree/quadtree
- k-d tree

Three Algorithms

- Barnes Hut
- Nearest Neighbor
- k Nearest Neighbor

Computation Patterns

- Aggregation (*sum*)
- Reduction (*e.g., min*)
- Sorting

Various Distance Metrics

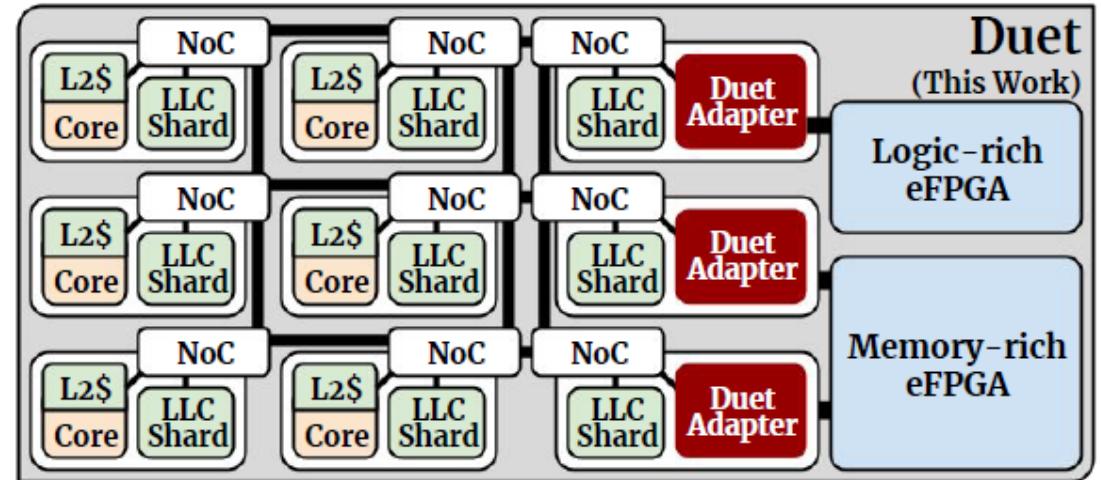
- Euclidean
- Manhattan
- Chebyshev



Evaluating an Open-Source SMHS: Duet

Duet

- A tightly-integrated, cache-coherent CPU-FPGA architecture
- Enables fine-grained transparent data sharing between the processors and the eFPGA-emulated accelerators
- Simulated in using **Gem5-Duet** extension



Duet[1]

| Platform | Backend | CPU | CPU Frequency | Accelerator | Accelerator frequency |
|--------------------------------|---------|-------------------------------|---------------|---------------|-----------------------|
| Duet (simulated in gem5) | HLS | RISC-V TimingSimple CPU | 1.5 GHz | Duet eFPGA | 333MHz |

Configuration



Grove Results Overview

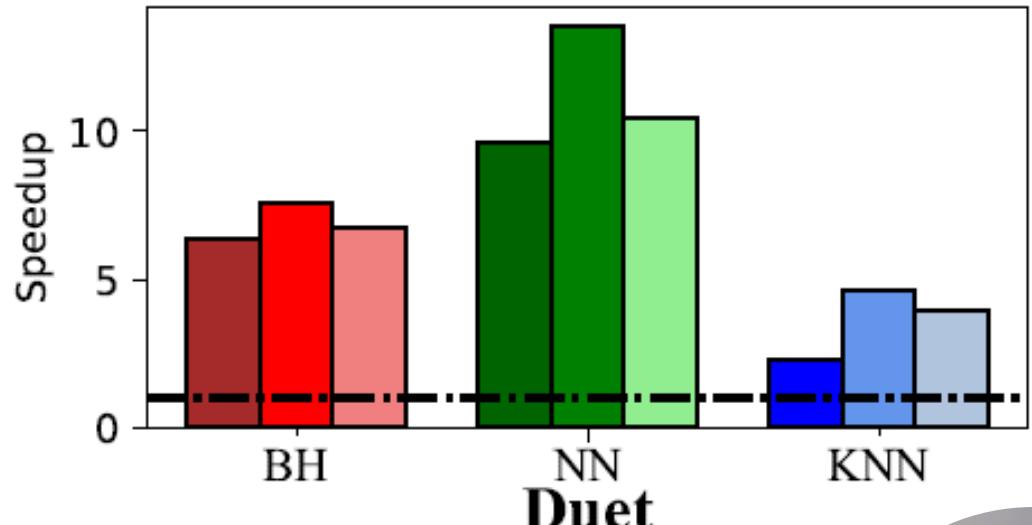
| | Leaf Size CPU | Leaf Size w/ Duet | Ratio | Avg Speedup |
|---------|---------------|-------------------|-------|-------------|
| BH | 3.33 | 512 | 153.6 | 6.9x |
| NN | 26.67 | 426.67 | 16 | 11.2x |
| KNN | 26.67 | 128 | 4.8 | 3.64x |
| Average | 19 | 355 | 18.8 | 6.43x |

We swept through the leaf node sizes to find the optimal configuration that yield the best performance,

- *Average 18x larger leaf node size than the CPU*

Speedups

highest **13.53x**
geomean **6.43x**



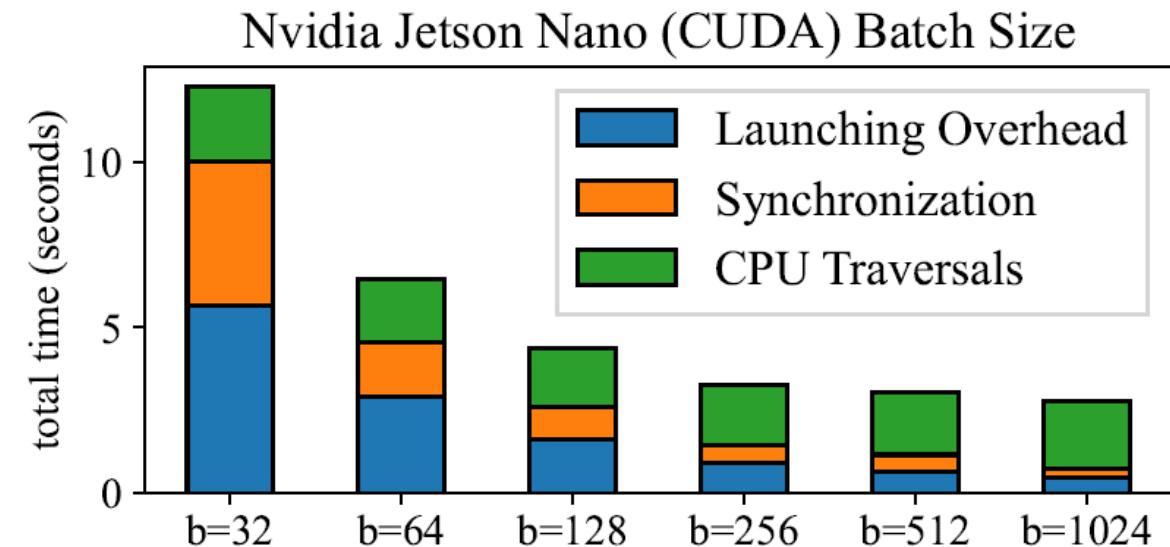
Speedups of the best heterogeneous configuration vs. the best homogeneous configuration of Grove.



We compared Duet to GPU-based SMHSs

Kernel Submission Cost

- Traverse-compute applications **frequently** invoke small kernels
- Useful works are shown in **Green**
- **Orange/Blue** are overheads
- Low-cost kernel submission is important for accelerating applications on edge devices
- **Duet has minimal offload overhead**

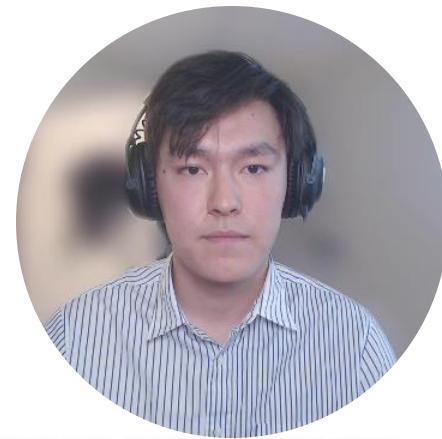


Batching multiple GPU kernels into a single/larger kernel helps amortizing kernel launching overhead on GPU-based systems



Conclusion

- ✓ We present how open-source hardware design can be used to accelerate a pragmatic class of applications
- ✓ We show that the Duet system can accelerate a suite of traverse-compute applications by up to **13.5x** with a geomean of **6.43x**
- ✓ We highlight the use of Grove, an open-source benchmark suite of traverse-compute workloads that utilize fine grained synchronization across PUs, and thus can provide a way for architecture researchers to evaluate their heterogeneous designs



UC Santa Cruz Redwood Grove

Team

Yanwen Xu yxu83@ucsc.edu 
Ang Li angl@princeton.edu
Tyler Sorensen tyler.sorensen@ucsc.edu

Open-Source Repo

Redwood & Grove at
<https://github.com/xuyanwen2012/redwood-rt>

